

Yohan Vergis Vinu

267-237-4755 | yohanvv@seas.upenn.edu | linkedin.com/in/yohanvinu | yohanvinu.com | github.com/YohanV1

EDUCATION

University of Pennsylvania <i>Master of Science in Engineering, Computer and Information Science</i>	Philadelphia, PA
• Relevant Coursework: Applied Machine Learning, Big Data Analytics, Internet and Web Systems	Aug. 2025 – May 2027
SRM Institute of Science and Technology <i>Bachelor of Technology, Computer Science and Engineering (GPA: 9.8/10)</i>	Chennai, India
• Leadership: Class Representative (8 semesters, Best CR award); IoT Lead/Founding Member, NetNex SRM	Aug. 2021 – May 2025

EXPERIENCE

Software Engineer Intern, Applied AI <i>Moatvie</i>	Aug. 2024 – June 2025
• Developed Vina , an agentic CHRO platform (featured on Product Hunt) that cuts recruiter screening time by 70% through interview question generation, performance analysis, and ranking (Django, PostgreSQL, LangGraph, AWS)	<i>Chennai, India</i>
• Built an AI permitting agent for Rowlett, TX, automating document verification and audit trails; reduced permit issuance from 3–4 weeks to minutes (Django, LangGraph, Tailwind CSS, AWS)	
• Designed a multi-channel utility agent for outages/billing, lowering call center workload by 60% via automated web, voice, and SMS support (Flask, LangGraph, Twilio, OpenAI Whisper, ElevenLabs, PostgreSQL)	
• Developed a Company Lookup API (FastAPI, ChromaDB, OpenAI SDK) achieving 95%+ NAICS accuracy using async scraping for real-time firmographics; and AI for CEQA compliance (Django, Gemini, ChromaDB, Leaflet.js) producing environmental reports in seconds and cutting review effort by 70%	
• Owned full product lifecycles: authored PRDs, shipped AI workflows, built CI/CD pipelines, and led client demos	
Machine Learning Engineer Intern <i>Samsung Research</i>	Aug. 2023 – Mar. 2024
• Fine-tuned and benchmarked ResNet , YOLOv8 , and MobileNetV2-SSD for multi-class pet detection, with MobileNetV2-SSD achieving a 7% accuracy gain over baseline methods	<i>Chennai, India</i>
• Compiled and annotated a 10K+ image dataset through manual image collection, labeling, and augmentation to improve class balance, earning Samsung's Certificate of Excellence	

PROJECTS

PennSearch: Distributed Cloud Search Engine <i>Java, AWS EC2, Distributed Systems</i>	Dec. 2025
• Built and deployed a fault-tolerant distributed web search engine with a custom web server, distributed key-value store, and parallel analytics engine	
• Implemented key-range sharding and replication with coordinator-based health checks to tolerate worker failures during distributed jobs	
• Engineered a crawl–index–rank pipeline indexing 300k+ pages with robots.txt compliance, language filtering, and multi-signal (TF-IDF, cosine similarity, PageRank) ranking achieving sub-50ms query latency	
Ensemble Models for Personal Expenditure Forecasting <i>PyTorch, LSTM, N-BEATS, LangGraph</i>	Apr. 2025
• Improved expenditure forecasting accuracy by 6% (MASE) by developing a custom ensemble over LSTM, N-BEATS, and dense networks	
• Built an agentic system to generate high-fidelity financial data spanning 24 attributes across 37 categories	
• Authored “Synthetic Data Generation and Neural Ensemble Architectures for Personal Expenditure Forecasting” (Accepted, IEEE INDICON 2025)	
InvoiceGPT: Agentic Invoice Processing <i>LangGraph, AWS (S3, RDS), TesseractOCR, Streamlit</i>	Dec. 2024
• Reduced manual invoice review time by 70% and achieved 95% OCR accuracy by building an AI-driven invoice processing platform using Tesseract/OpenAI on AWS	
• Reduced query resolution from minutes to seconds by developing a LangGraph-powered conversational and visual agent for natural language financial analytics	

TECHNICAL SKILLS

Programming Languages: Python, SQL, Java, TypeScript, JavaScript, C/C++
AI/ML: PyTorch, TensorFlow, NumPy, Pandas, OpenCV, Hugging Face, LangGraph/LangChain, LangSmith
Web Development & APIs: React, Next.js, Django, Flask, FastAPI, Streamlit, Tailwind CSS, REST APIs, Postman
Cloud, Databases & DevOps: AWS, GCP, Azure, PostgreSQL, ChromaDB, Docker, Git/GitHub, GitHub Actions